

Recurrent Generative Stochastic Networks for Sequence Prediction

Abstract

We present a new generative model for unsupervised learning of sequence representations, the recurrent Generative Stochastic Network (RNN-GSN). This approach builds on recent deep recurrent neural network architectures for modeling high-dimensional, multimodal input distributions, and is able to predict complex sequences of these distributions. The RNN-GSN is non-probabilistic and easier to perform inference over compared to other representation models, such as the RNN-RBM or NADE variants. We provide experimental validation for this new model on sequences of MNIST images and standard MIDI music datasets.

1. Introduction

Unsupervised sequence learning is an important problem in machine learning given that most information (ranging from speech to video to music to even consumer behavior) is often unlabeled and has a sequential structure. Most of these sequences consist of high-dimensional, complex objects such as words in text, images in video, or chords in music. Recently, recurrent neural networks (RNN) (Rumelhart et al., 1986) have become state-of-the-art for sequence representation because they have an internal memory that can learn long-term, temporal dependencies. Further advances with Hessian-free optimization and Long Short-term Memory for RNNs (Martens & Sutskever, 2011; Hochreiter & Schmidhuber, 1996) have enabled learning highly complex temporal dependencies.

While traditional RNNs can learn complex sequences through maintaining an internal memory, they don't deal well when modeling complex input distributions, where the conditional distribution at each time step is highly multimodal. In real-world data, we often care more about this conditional distribution rather than the expected value. In music, for example, the existence of a particular note can greatly change the probabilities with which other notes occur at the same time in a chord. In consumer behavior, the

existence of a particular action taken can greatly change the probabilities of other actions made within the same time window. We would prefer to reason and perform inference over these distributions at each time step to form a better representation of the sequences.

Deep RNN architectures have been proposed to alleviate the traditional RNN architecture shortcomings with complex input distributions (Pascanu et al., 2013). One proposed framework involves using deep input-to-hidden or hidden-to-output functions to reduce input complexity, making the RNN's temporal modeling task easier. These functions exploit the ability of deep networks to disentangle the underlying factors of variation of the original input and flatten high-density data manifolds (Goodfellow et al., 2009; Glorot et al., 2011; Bengio et al., 2012).

One such model, the RNN-RBM, replaces the output layer of an RNN with a restricted Boltzmann machine (RBM) to provide a conditional generative model over the input distribution (Boulanger-Lewandowski et al., 2012; Boulanger-Lewandowski, 2014). The RNN-RBM is a generalization of previous recurrent temporal RBM work (Sutskever et al., 2009) that has shown promise for modeling complex sequences such as MIDI representations of polyphonic music. Another possible model is to replace the RNN output layer with a neural autoregressive distribution estimator (NADE), which has been shown to provide a tractable distribution estimator that performs similarly to large, intractable RBMs (Larochelle & Murray, 2011).

Motivated by the promise of deep RNN models for learning sequence representations, we propose a model using a Generative Stochastic Network (GSN) as the hidden-to-output function of an RNN, called the RNN-GSN. GSNs are a recent generalization of denoising autoencoders (Bengio et al., 2013a) that are easier to perform inference over compared to RBMs and sample from compared to NADEs. GSNs are non-probabilistic, generative models that have also been shown to learn the same model as a deep orderless NADE, which alleviates the factorization ordering drawback of a traditional NADE model (Yao et al., 2014).

The rest of this paper explains the RNN-GSN model and provides experimental evidence for its use in modeling complex sequences. In Sections 2, 3, and 4, we introduce the GSN, RNN, and RNN-GSN architectures in more detail. We then provide experimental validation of the RNN-

GSN on sequences of MNIST images and standard MIDI datasets in Section 5.

2. Generative Stochastic Networks

Generative stochastic networks (GSN) are a generalization of the denoising auto-encoder and help solve the problem of mixing between many major modes of the input data distribution.

2.1. Denoising auto-encoder

Denoising auto-encoders use a Markov chain to learn a reconstruction distribution $P(X|\tilde{X})$ given a corruption process $C(\tilde{X}|X)$ for some data X . Denoising auto-encoders have been shown as generative models (Bengio et al., 2013b), where the Markov chain can be iteratively sampled from:

$$X_t \sim P_{\Theta}(X|\tilde{X}_{t-1}) \quad (1)$$

$$\tilde{X}_t \sim C(\tilde{X}|X_t) \quad (2)$$

As long as the learned distribution $P_{\Theta_n}(X|\tilde{X})$ is a consistent estimator of the true conditional distribution $P(X|\tilde{X})$ and the Markov chain is ergodic, then as $n \rightarrow \infty$, the asymptotic distribution $\pi_n(X)$ of the generated samples from the denoising auto-encoder converges to the data-generating distribution $P(X)$ (Bengio et al., 2013b).

2.2. Easing restrictive conditions on the denoising auto-encoder

A few restrictive conditions are necessary to guarantee ergodicity of the Markov chain - requiring $C(\tilde{X}|X) > 0$ everywhere that $P(X) > 0$. Particularly, a large region V containing any possible X is defined such that the probability of moving between any two points in a single jump $C(\tilde{X}|X)$ must be greater than 0. This restriction requires that $P_{\Theta_n}(X|\tilde{X})$ has the ability to model every mode of $P(X)$, which is a problem this model was meant to avoid.

To ease this restriction, Bengio et al. (Bengio et al., 2013a) prove that using a $C(\tilde{X}|X)$ that only makes small jumps allows $P_{\Theta}(X|\tilde{X})$ to model a small part of the space V around each \tilde{X} . This weaker condition means that modeling the reconstruction distribution $P(X|\tilde{X})$ would be easier since it would probably have fewer modes.

However, the jump size σ between points must still be large enough to guarantee that one can jump often enough between the major modes of $P(X)$ to overcome the regions of low probability: σ must be larger than half the largest distance of low probability between two nearby modes, such that V has at least a single connected component between modes. This presents a tradeoff between the difficulty of learning $P_{\Theta}(X|\tilde{X})$ and the ease of mixing between modes

separated by this low probability region.

2.3. Generalizing to GSN

While denoising auto-encoders can rely on X_t alone through a deterministic procedure for the state of the Markov chain, GSNs introduce a latent variable H_t that acts as an additional state variable in the Markov chain along with the visible X_t (Bengio et al., 2013a):

$$H_{t+1} \sim P_{\Theta_1}(H|H_t, X_t) \quad (3)$$

$$X_{t+1} \sim P_{\Theta_2}(X|H_{t+1}) \quad (4)$$

The latent state variable H can be equivalently defined as $H_{t+1} = f_{\Theta_1}(X_t, Z_t, H_t)$, a learned function f with an independent noise source Z_t such that X_t cannot be reconstructed exactly from H_{t+1} . If X_t could be recovered from H_{t+1} , the reconstruction distribution would simply converge to the Dirac at X . Denoising auto-encoders are therefore a special case of GSNs, where f is fixed instead of learned.

GSNs also use the notion of walkbacks to aid training. Walkbacks are the process of generating samples by iteratively sampling from $P_{\Theta_1}(H|H_t, X_t)$ and $P_{\Theta_2}(X|H_{t+1})$ for a given input. By using walkbacks, the model is more likely to seek out spurious modes in the data distribution and correct for them (Bengio et al., 2013b). The resulting Markov chain of a GSN is inspired by Gibbs sampling, but with stochastic units at each layer that can be backpropagated (Rezende et al., 2014).

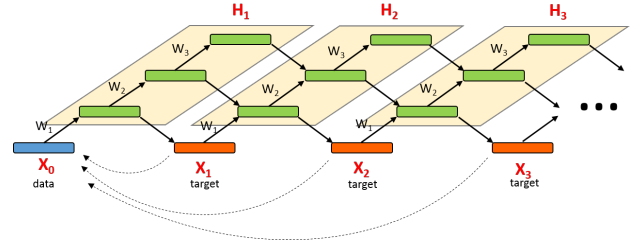


Figure 1. Unrolled GSN Markov chain.

Experimental results with GSNs show that their latent states mix well between the major modes of the data - mixing faster at higher layers in the model (Bengio et al., 2013a). Using this property, we tested a simple temporal GSN model to predict sequences of inputs. The temporal GSN uses a linear transformation $H \rightarrow H$ to predict $P(H_t|H_{t-1}, \dots, H_{t-n})$ with an input history of size n . Using this predicted H_t , an expected input x_t was sampled from the model's learned reconstruction distribution $P_{\Theta_2}(X|H)$.

Qualitatively, this model appears to predict temporal dependencies well within its history window n (Figure 3).

This result provided motivation to use the latent state H as the emitted parameter in a better-suited temporal model, such as an RNN.




Figure 2. Temporal GSN with history $n = 2$ trained on an arbitrary MNIST sequence. Original input sequence is on the left, a noisy version of the input fed into the model is in the middle, and the predicted input based on the history is on the right.

3. Recurrent Neural Networks

Traditional RNNs simulate a discrete-time system that has an input sequence $\{x_1 \dots x_t\}$ and a hidden state sequence $\{h_1 \dots h_t\}$ that map to an output sequence $\{y_1 \dots y_t\}$. The network is defined for timestep t in the sequence by the hidden and output equations:

$$h_t = \Phi_h(W^T h_{t-1} + U^T x_t + b_h) \quad (5)$$

$$y_t = \Phi_y(V^T h_t + b_y) \quad (6)$$

where Φ_h and Φ_y are element-wise nonlinear functions and W , U , V , b_h , and b_y are the network parameters Θ . The hidden state h contains all information about the sequence for a given timestep t .

RNNs can be trained via backpropagation over a set number of time steps, which is known as backpropagation through time. However, due to the exploding gradient problem, backpropagation through time has difficulty training the parameters for long-term temporal dependencies. To fix the exploding gradient problem, long short-term memory units (LSTM) (Hochreiter & Schmidhuber, 1996) can be introduced as the hidden units that contain extra parameters to learn a gated memory for storing the unit's activation value for a period of time steps. With this architecture, normal backpropagation through time can train the model to learn much longer temporal dependencies. Alternatively, Hessian-free optimization (Martens & Sutskever, 2011) can learn the model parameters for long-term temporal dependencies without needing to introduce specialized hidden units.

3.1. Extension to deep RNN

4. The RNN-GSN

The Recurrent GSN (RNN-GSN)

5. Experiments

This is the experiments section.

5.1. Sequences of MNIST digits

arbitrary sequences of images.

- Sequence1 is a simple linear sequence of digits 0-9 repeating.
- Sequence2 introduces one bit of parity by alternating sequences 0-9 and 9-0 repeating, where the next value depends on whether the sequence is ascending or descending.
- Sequence3 creates a longer, more complex sequence by introducing a second bit of parity. It is formed by:

5.2. Sequences of polyphonic music

midi stuff.

- Piano-midi.de
- Nottingham
- MuseData
- JSB chorales

6. Conclusion

This is the conclusion.

References

- Bengio, Yoshua, Mesnil, Grégoire, Dauphin, Yann, and Rifai, Salah. Better mixing via deep representations. *CoRR*, abs/1207.4404, 2012.
- Bengio, Yoshua, Thibodeau-Laufer, Eric, and Yosinski, Jason. Deep generative stochastic networks trainable by backprop. *CoRR*, abs/1306.1091, 2013a.
- Bengio, Yoshua, Yao, Li, Alain, Guillaume, and Vincent, Pascal. Generalized denoising auto-encoders as generative models. *CoRR*, abs/1305.6663, 2013b.
- Boulanger-Lewandowski, Nicolas. *Modeling High-Dimensional Audio Sequences with Recurrent Neural Networks*. PhD thesis, Université de Montréal, april 2014.
- Boulanger-Lewandowski, Nicolas, Bengio, Yoshua, and Vincent, Pascal. Modeling temporal dependencies in high-dimensional sequences: Application to polyphonic music generation and transcription. In *Proceedings of the 29th International Conference on Machine Learning*, 2012.

330	Glorot, Xavier, Bordes, Antoine, and Bengio, Yoshua.	385
331	Domain adaptation for large-scale sentiment classifica-	386
332	tion: A deep learning approach. In <i>In Proceedings of</i>	387
333	<i>the Twenty-eight International Conference on Machine</i>	388
334	<i>Learning, ICML, 2011.</i>	389
335		390
336	Goodfellow, Ian, Lee, Honglak, Le, Quoc V., Saxe, An-	391
337	drew, and Ng, Andrew Y. Measuring invariances in	392
338	deep networks. In Bengio, Y., Schuurmans, D., Lafferty,	393
339	J.D., Williams, C.K.I., and Culotta, A. (eds.), <i>Advances</i>	394
340	<i>in Neural Information Processing Systems 22</i> , pp. 646–	395
341	654. Curran Associates, Inc., 2009.	396
342		397
343	Hochreiter, Sepp and Schmidhuber, Jürgen. Bridging long	398
344	time lags by weight guessing and "long short term mem-	399
345	ory". In <i>Spatiotemporal models in biological and artifi-</i>	400
346	<i>cial systems</i> , pp. 65–72. IOS Press, 1996.	401
347		402
348	Larochelle, Hugo and Murray, Iain. The neural autoregres-	403
349	sive distribution estimator. <i>Journal of Machine Learning</i>	404
350	<i>Research</i> , 15:29–37, 2011.	405
351		406
352	Martens, James and Sutskever, Ilya. Learning recurrent	407
353	neural networks with hessian-free optimization. In <i>Pro-</i>	408
354	<i>ceedings of the 28th International Conference on Ma-</i>	409
355	<i>chine Learning (ICML-11)</i> , pp. 1033–1040, 2011.	410
356		411
357	Pascanu, Razvan, Gülçehre, Çağlar, Cho, Kyunghyun, and	412
358	Bengio, Yoshua. How to construct deep recurrent neural	413
359	networks. <i>CoRR</i> , abs/1312.6026, 2013.	414
360		415
361	Rezende, Danilo J., Mohamed, Shakir, and Wierstra, Daan.	416
362	Stochastic backpropagation and approximate inference	417
363	in deep generative models. In <i>Proceedings of the 31st</i>	418
364	<i>International Conference on Machine Learning, 2014.</i>	419
365		420
366	Rumelhart, David E., Hinton, Geoffrey E., and Williams,	421
367	Ronald J. Learning representations by back-propagating	422
368	errors. In <i>Parallel Dist. Proc.</i> , pp. 318–362. MIT Press,	423
369	1986.	424
370		425
371	Sutskever, Ilya, Hinton, Geoffrey E., and Taylor, Gra-	426
372	ham W. The recurrent temporal restricted boltzmann	427
373	machine. In Koller, D., Schuurmans, D., Bengio, Y., and	428
374	Bottou, L. (eds.), <i>Advances in Neural Information Pro-</i>	429
375	<i>cessing Systems 21</i> , pp. 1601–1608. Curran Associates,	430
376	Inc., 2009.	431
377		432
378	Yao, Li, Ozair, Sherjil, Cho, Kyunghyun, and Bengio,	433
379	Yoshua. On the equivalence between deep nade and	434
380	generative stochastic networks. <i>CoRR</i> , abs/1409.0585,	435
381	2014.	436
382		437
383		438
384		439